

Building PyTorch Models for Image Classification and Reconstruction



Dr. Susant Kumar Panigrahi

SERB Sponsored N-PDF

Department Of Electrical Engineering

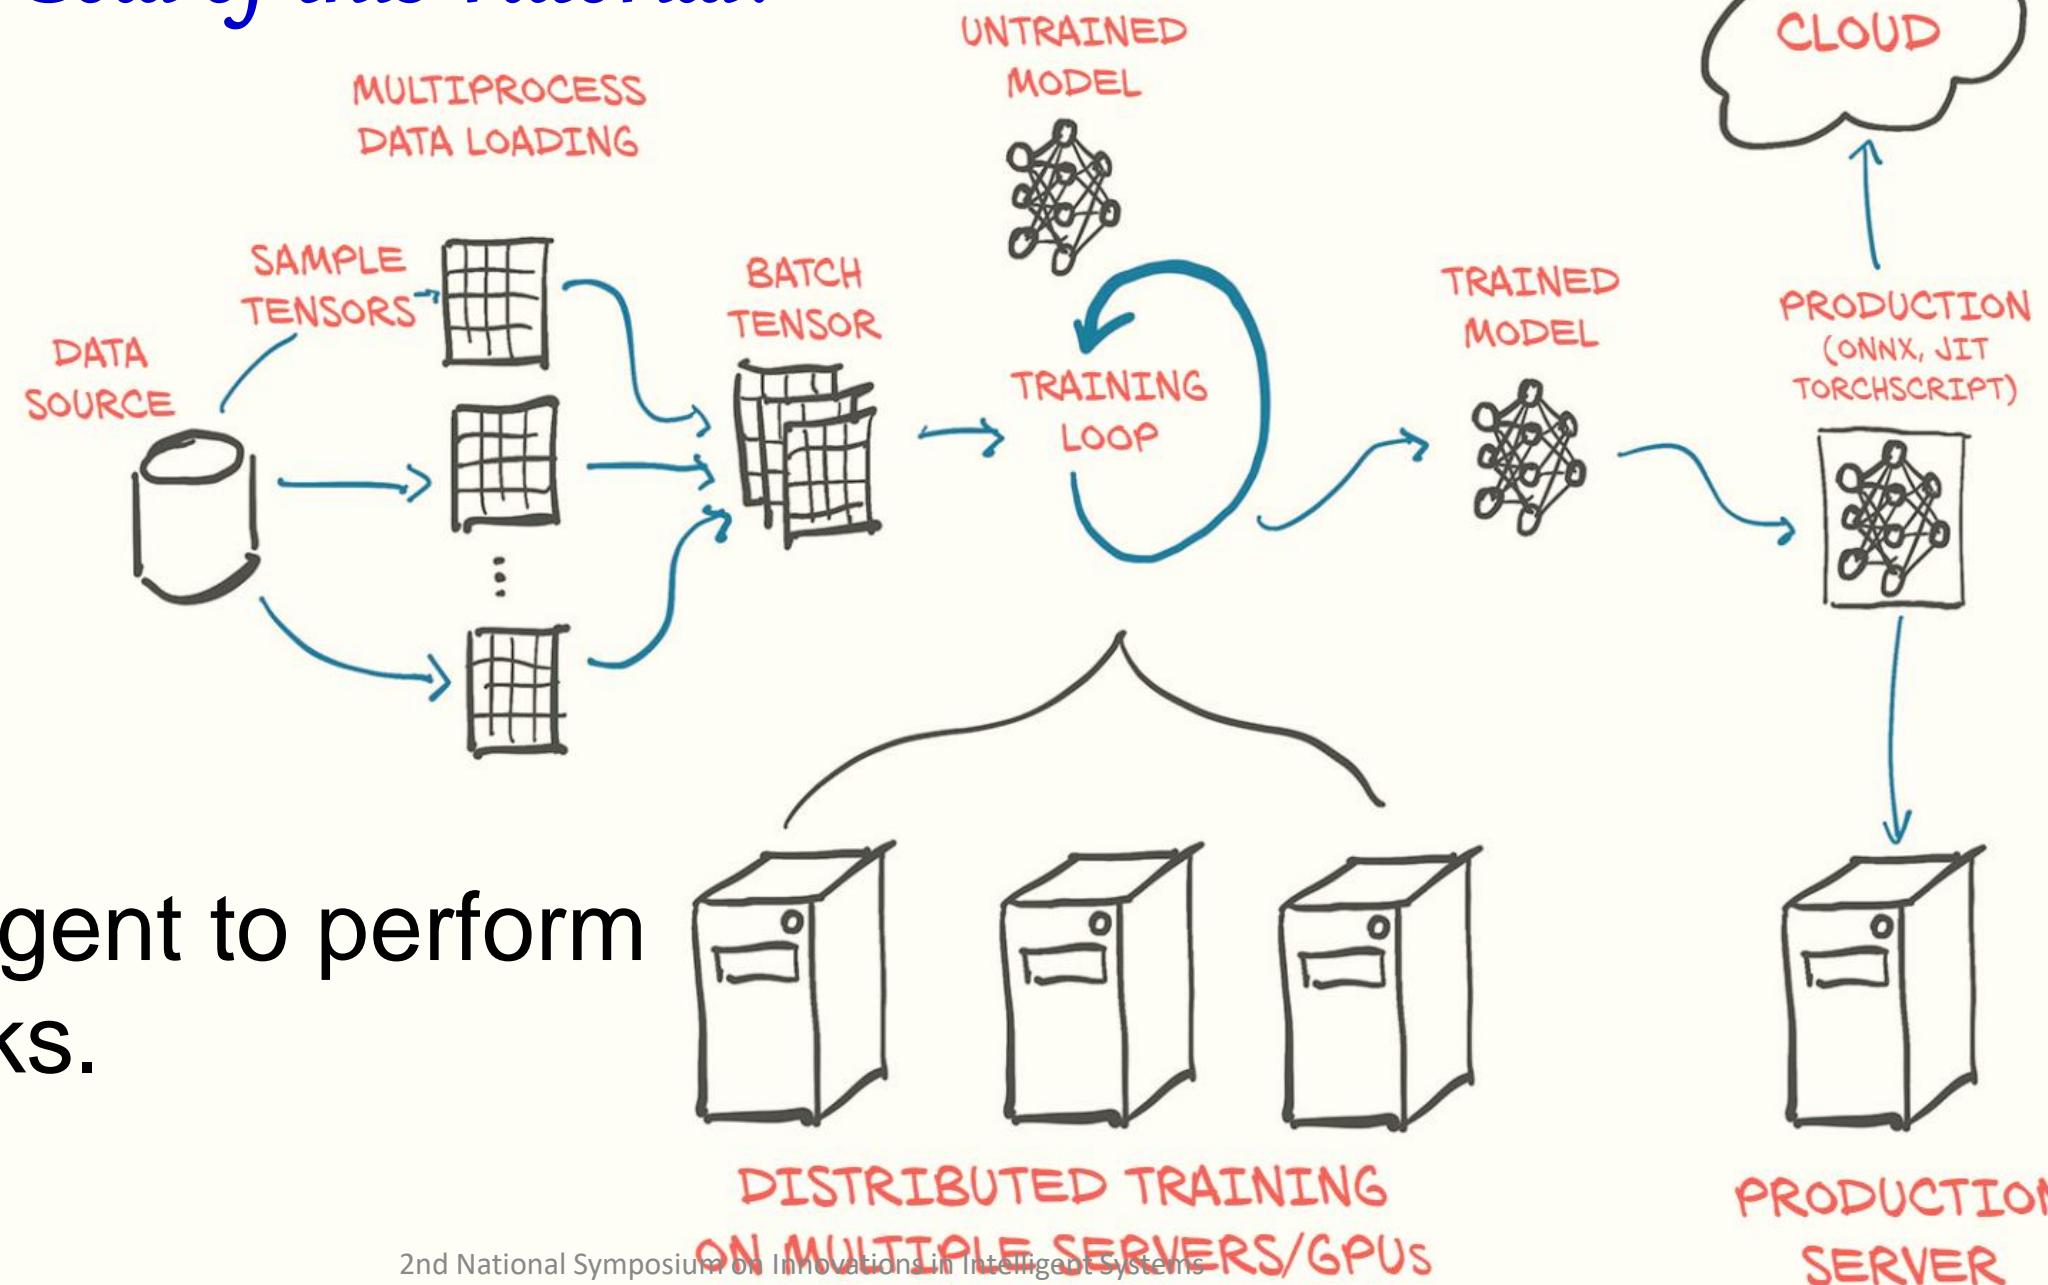
Indian Institute of Technology Kharagpur

Email: spanigrahi@kgipian.iitkgp.ac.in



PyTorch

What is the Goal of this Tutorial?



What is the Goal of this Tutorial?

How do train a model?

$$w^* = \underset{w}{\operatorname{argmin}}$$

Gradient
Descent

$$\sum_{(x,y) \in D} \mathcal{L}(f_w(x), y)$$

Dataset

Loss Function

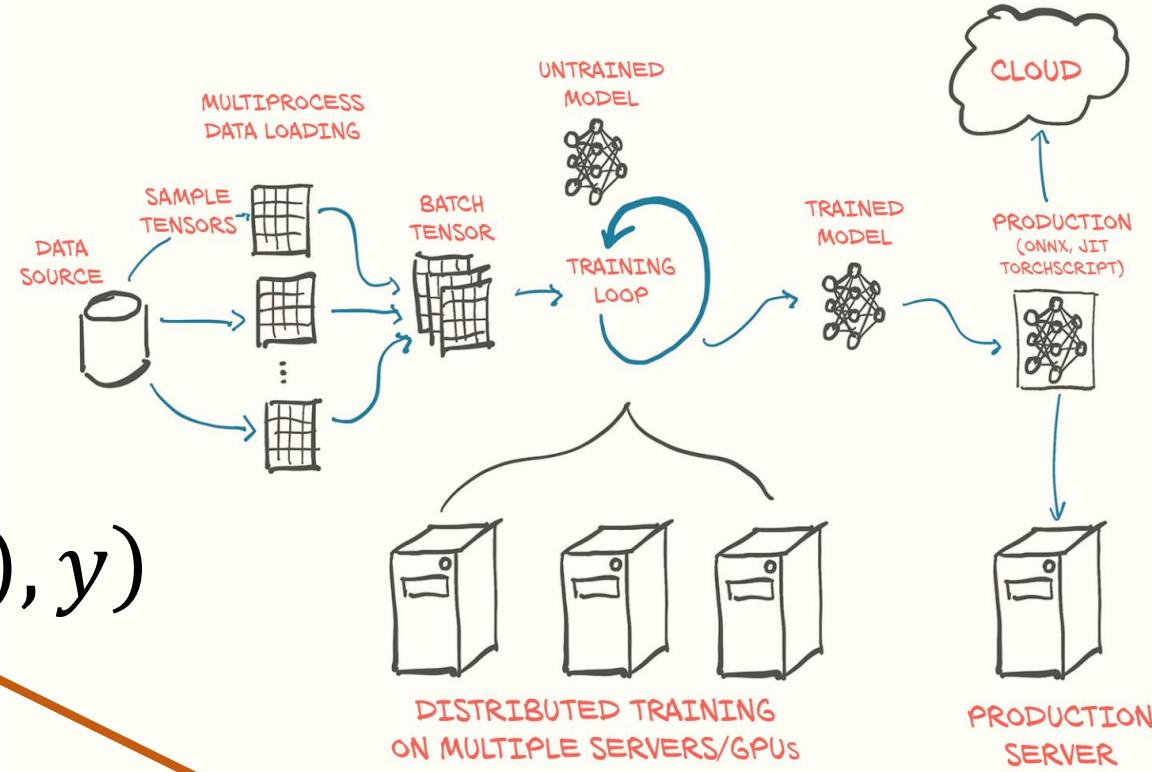
Model

Neural Network or
Convolution Neural Network

PyTorch does all of these!

20 February 2025

2nd National Symposium on Innovations in Intelligent Systems



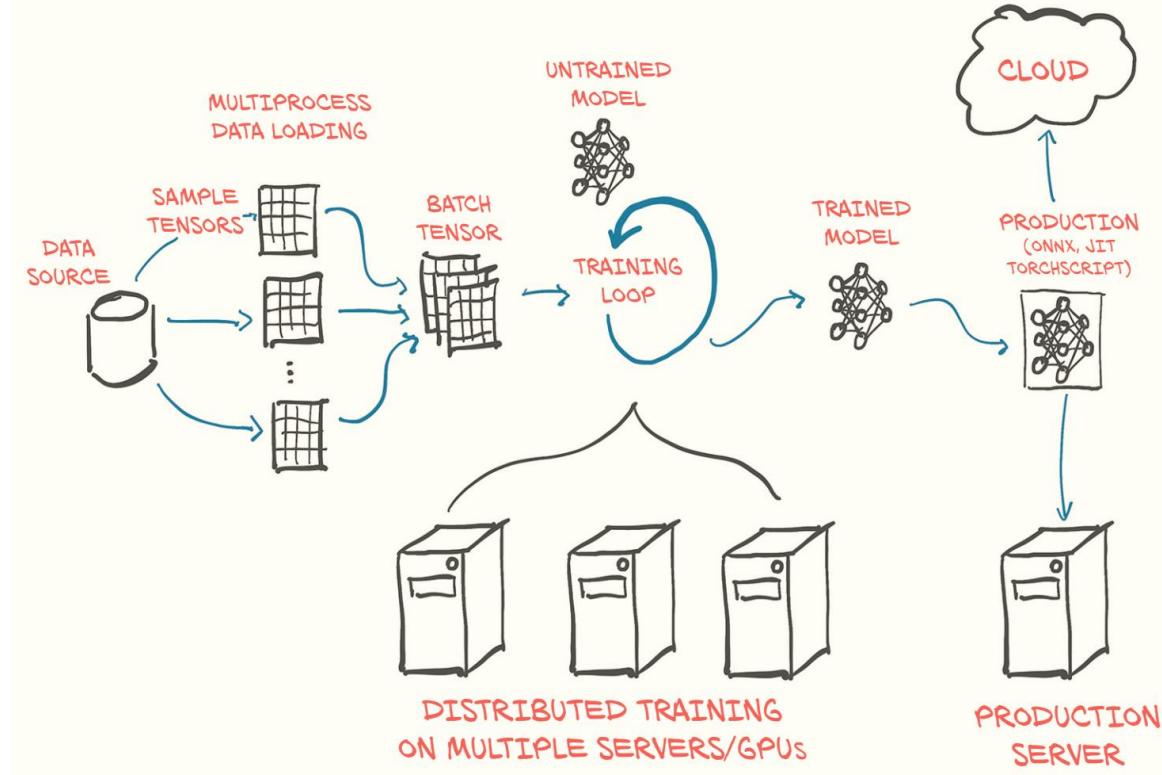
What is PyTorch?

Python Library for:

- Defining Neural Networks
- Automatic computation of gradients
- Network training and model inference
- And more ! (dataloader, optimizer, scheduler, Cuda and Many more.)



2nd National Symposium on Innovations in Intelligent Systems



$$w^* = \underset{w}{\operatorname{argmin}} \sum_{(x,y) \in D} \mathcal{L}(f_w(x), y)$$

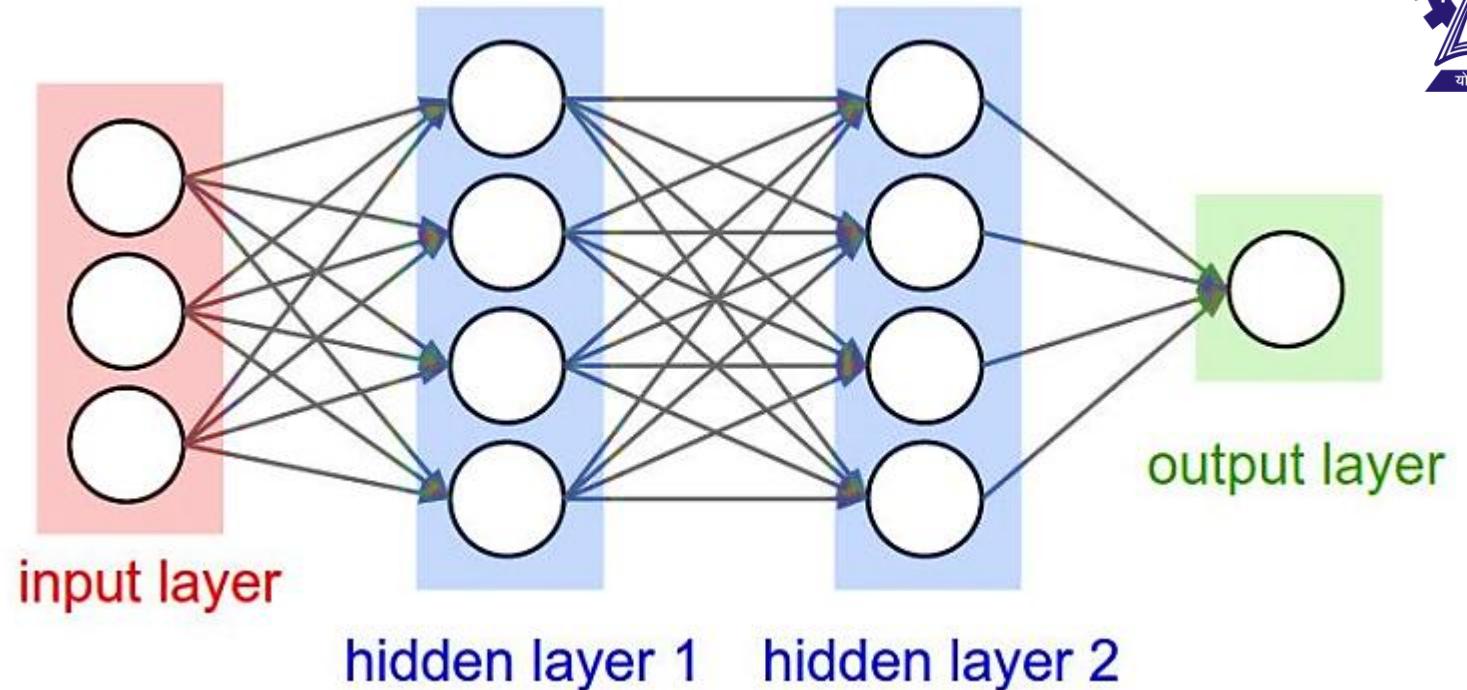
Diagram illustrating the optimization process:

- The expression $w^* = \underset{w}{\operatorname{argmin}} \sum_{(x,y) \in D} \mathcal{L}(f_w(x), y)$ is shown.
- The term $\underset{w}{\operatorname{argmin}}$ is highlighted with a green oval.
- The term $\sum_{(x,y) \in D}$ is highlighted with a blue oval.
- The term $\mathcal{L}(f_w(x), y)$ is highlighted with an orange oval.
- An arrow labeled 'Dataset' points to the summation term.
- An arrow labeled 'Loss Function' points to the highlighted loss term.
- An arrow labeled 'Network' points to the highlighted function term.
- An arrow labeled 'Gradient Descent' points from the highlighted argmin term to the bottom left.

How does PyTorch Work?



PyTorch and
Computational Graph



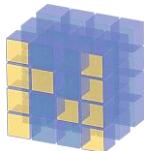
Definition:

$$h_1 = \sigma(W_1x) \quad h_2 = \sigma(W_2h_1) \quad y = \sigma(W_3h_2)$$

PyTorch Computation:

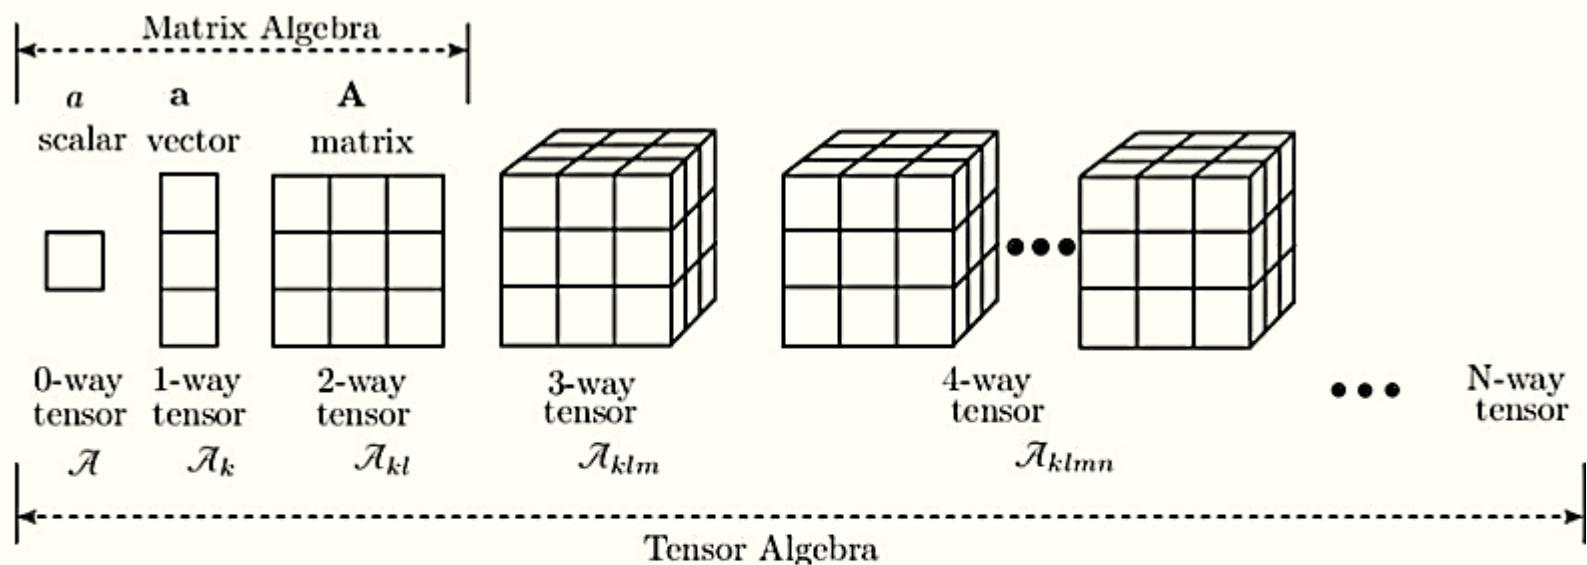
$$\frac{\partial y}{\partial W_1} = \frac{\partial y}{\partial h_2} \frac{\partial h_2}{\partial h_1} \frac{\partial h_1}{\partial W_1} \quad \frac{\partial y}{\partial W_2} = \frac{\partial y}{\partial h_2} \frac{\partial h_2}{\partial W_2} \quad \frac{\partial y}{\partial W_3}$$

Multidimensional Arrays Computation



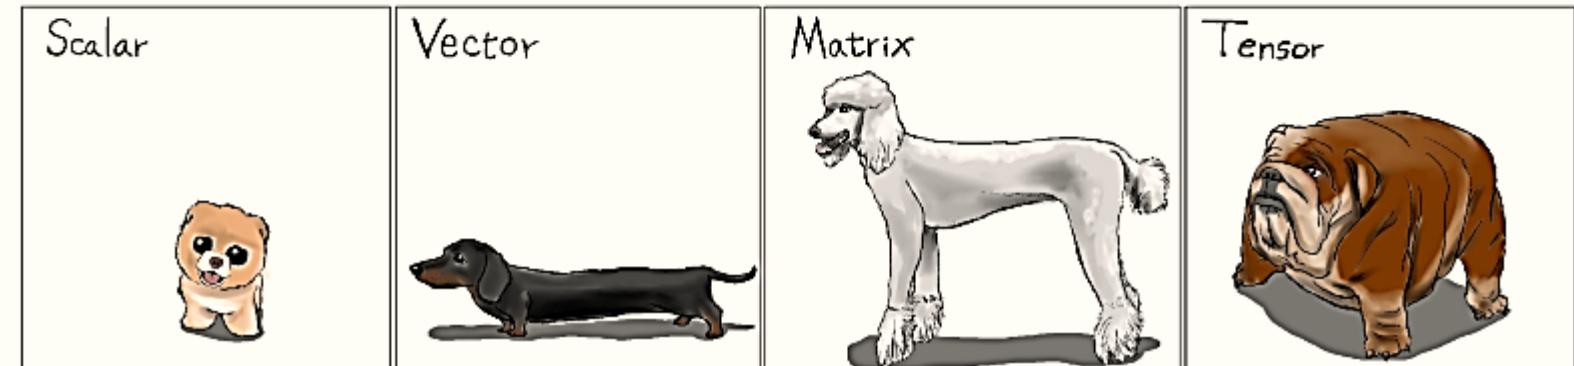
NumPy

- ✓ Fast CPU implementations
- ✓ CPU-only
- ✓ No autodiff
- ✓ Imperative



PyTorch

- ✓ Fast CPU implementations
- ✓ Allows GPU
- ✓ Supports autodiff
- ✓ Imperative



Other features include:

- Datasets and dataloading
- Common neural network operations
- Built-in optimizers (Adam, SGD, ...)



PyTorch Basics

```
import torch
import numpy as np

# Creating an empty tensor.

x = torch.empty(3,3)

x1 = torch.ones(3,3)

x2 = torch.rand(3,3)

x3 = torch.zeros(3,3)

print(x)

print('Torch ones matrix:', x1)

print('Torch rand matrix:', x2)

print('Torch zeros matrix:', x3)
```

2. Basic Operations with Torch tensor

```
x = torch.rand(2,2)
y = torch.rand(2,2)
z = x + y print('Addition:', z)
z1 = torch.add(x,y) print('Addition:', z1)
# Inplace addition
y.add_(x)
print('Inplace Addition:', y)

# Other Operations
z = y - x
print('Subtraction:', z)
z1 = torch.subtract(y,x)
print(f'Subtraction: {z1}')
z1.subtract_(y)
print(f'Inplace subtraction: {z1}')
```

https://github.com/susant146/PyTorch_Basics_CNNmodels/blob/main/SI_01_Pytorch_Basics.ipynb



PyTorch Basics

Indian Institute of Technology Kharagpur

Department of Electrical Engineering | Signal Processing and Machine Learning

```
x = torch.rand(2,2)
y = torch.rand(2,2)
z = x*y
print('Multiplication: ', z)
z1 = torch.multiply(x,y) # torch.mul and torch.multiply
print('Element-Wise Multiplication:', z1)
```

```
# Inplace multiplication: y.multiply_(x)
z = x / y # torch.div(x,y)
print('Element-Wise Division: \n', z)
y.div_(x)
print('Inplace Division: \n', y) # Matrix multiplications
```

```
# Define two matrices
A = torch.tensor([[1, 2], [3, 4]]) # 2x2 matrix
B = torch.tensor([[5, 6, 1], [7, 8, 0]]) # 2x3 matrix
```

3. Tensor Multiplication and Division

```
# Alternative method: Using torch.matmul (supports broadcasting for higher dimensions)
C_broadcast = torch.matmul(A, B)
print("\nMatrix Multiplication Result (torch.matmul):") print(C_broadcast)
```



PyTorch Basics

Indian Institute of Technology Kharagpur

Department of Electrical Engineering | Signal Processing and Machine Learning

```
# Create a tensor
```

```
tensor = torch.arange(12)
print("Original Tensor:")
print(tensor)
```

```
# Reshape to 3x4 matrix
```

```
reshaped_tensor = tensor.reshape(3, 4)
print("\nReshaped to 3x4:")
print(reshaped_tensor)
```

```
flattened_tensor = reshaped_tensor.reshape(-1)
print("\nFlattened Tensor:")
print(flattened_tensor)
```

4. Tensor Reshape & View

```
# Reshape to 2x6 using view # reshape:  
Works regardless of tensor's memory layout.
```

```
# view: Faster but requires the tensor to  
have contiguous memory. Use .contiguous()  
before view if needed.
```

```
reshaped_view = tensor.view(2, 6)
print("\nReshaped to 2x6 using view:")
print(reshaped_view)
```

```
# Ensuring contiguity before view
```

```
non_contiguous = tensor.t()
contiguous_tensor =
non_contiguous.contiguous().view(-1)
```

```
# Automatically infer one dimension # Use -1 to let PyTorch infer the size of one dimension.
```

```
inferred_shape = tensor.reshape(3, -1)
print("\nReshaped with Inferred Dimension (2x-1):") print(inferred_shape)
```



PyTorch Basics

Indian Institute of Technology Kharagpur

Department of Electrical Engineering | Signal Processing and Machine Learning

5. Numpy to PyTorch Tensor & Vice-Versa

```
a = torch.ones(5)
print(a)
b = a.numpy()
print('Numpy array: \n', b)
print('Type: \t', type(b))

# need to be very careful with place
operations
a.add_(2)
print('Modified a: \n', a)
print('See What happend! \n', b)
```

```
a = np.ones(5) print(a) b =
torch.from_numpy(a)

# default dtype is float64, change
# b = b.to(dtype=torch.int)
# or b = b.type(torch.IntTensor)
print(b)
# Inplace operation
a += 1
print('a ndarray: \t',a)
print('b Tensor: \t',b)
# Check this one
a = a+1
print('a ndarray: \t',a)
print('b Tensor: \t',b)
```



PyTorch Basics

6. Check for Available Device: Cuda or CPU

```
if torch.cuda.is_available():
    device = torch.device("cuda") # Use GPU
    print("CUDA is available. Using GPU.")
else:
    device = torch.device("cpu") # Use CPU
    print("CUDA is not available. Using CPU.")
```

```
x = torch.ones(5, device=device)
y = torch.ones(5)
y = y.to(device)
z = x + y
```

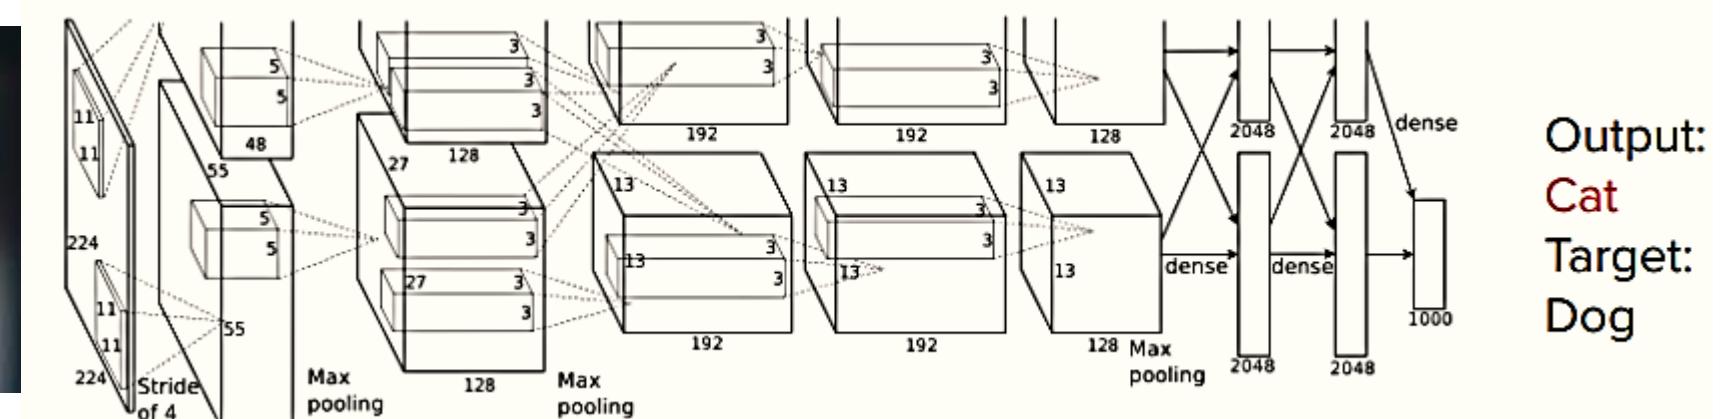
```
print(f'The operation performed in {device}, output: {z}')
# As x, y and z are in CUDA, we can't convert them to numpy
# z1 = z.numpy() # Error Here
# bringing z to cpu
z1 = z.cpu() z1 = z1.numpy()
```

```
print(f"z1 is in cpu, output ndarray: {z1}")
```

Move a Tensors to CUDA

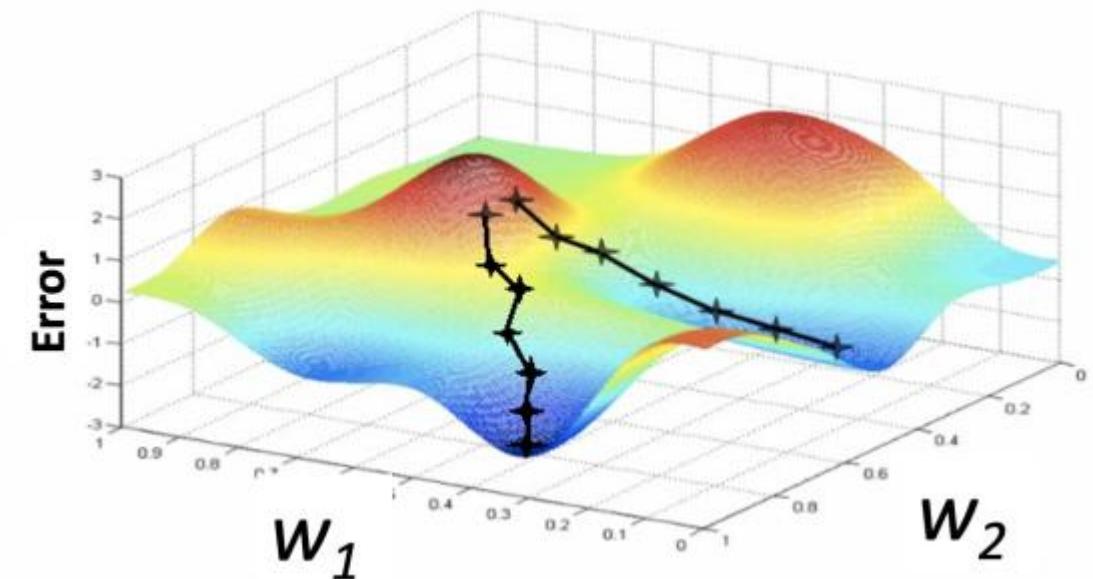
Gradient Descent & Backward Pass

Computational graph and Autograd with Pytorch



Learning as optimization:

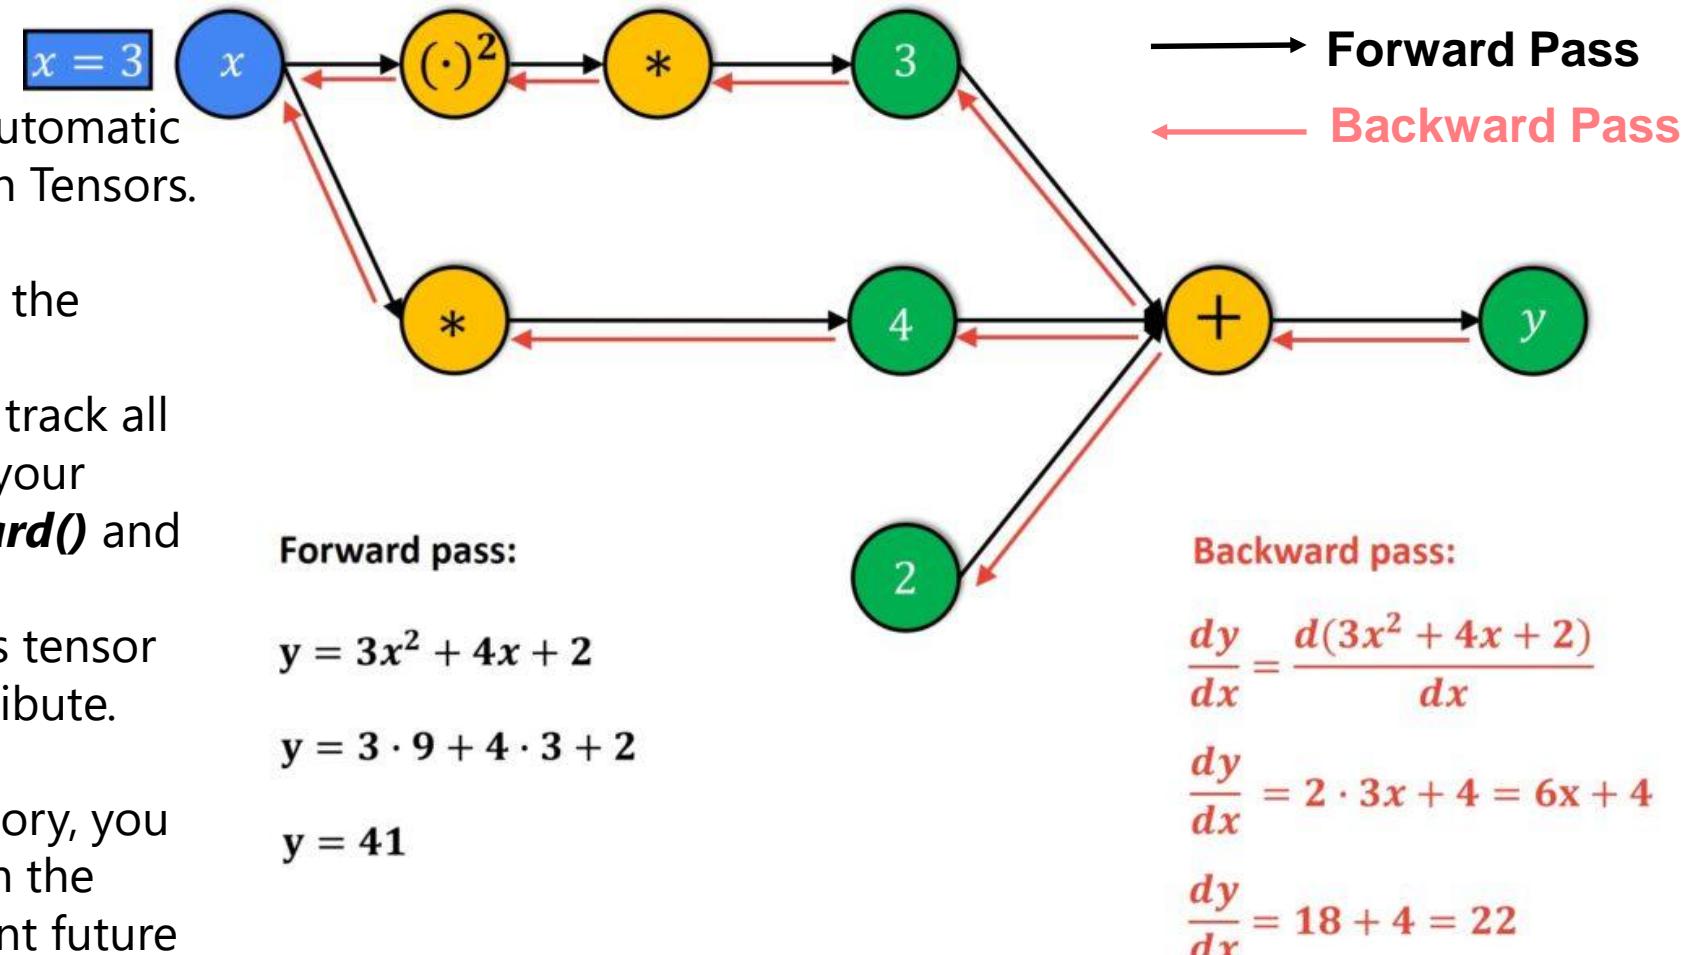
1. An input-output function (an ANN): $y = f_w(x)$
2. A Loss function: $L = \mathcal{L}(y, y_{ref})$
3. Optimization problem: $w^* = \operatorname{argmin}_w \sum_{(x, y_{ref}) \in D} \mathcal{L}(f_w(x), y_{ref})$



Gradient Descent & Backward Pass

Computational graph and Autograd with Pytorch

- ✓ The ***autograd*** package provides automatic differentiation for all operations on Tensors.
- ✓ ***torch.Tensor*** is the central class of the package. If you set its attribute ***.requires_grad*** as True, it starts to track all operations on it. When you finish your computation you can call ***.backward()*** and have all the gradients computed automatically. The gradient for this tensor will be accumulated into ***.grad*** attribute.
- ✓ To stop a tensor from tracking history, you can call ***.detach()*** to detach it from the computation history, and to prevent future computation from being tracked.



Gradient Descent & Backward Pass

Computational graph and Autograd with Pytorch

1. Autograd

```
import torch

x = torch.tensor(3., requires_grad=True)
y = 3*x**2 + 4*x + 2
print(y)

y.backward()
print(x.grad)
```

2. How to turn off the gradient calculations?

```
x = torch.tensor(3., requires_grad=True)
print(x)

x = x.requires_grad_(False)
print(x)

x = x.detach()
print(x)
```

3. Gradient accumulation effect

```
x = torch.tensor(3., requires_grad=True)

for epoch in range(3):
    y = 3*x**2 + 4*x + 2
    y.backward()

    print(x.grad)
```

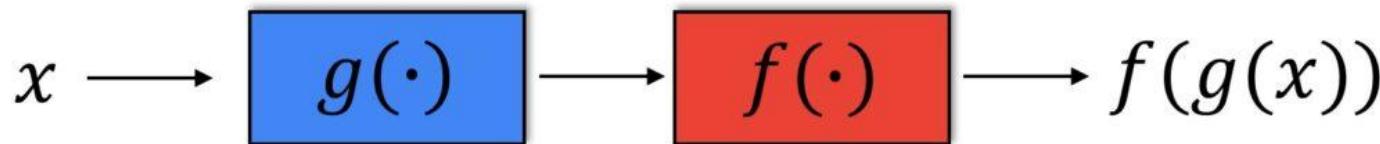
```
x = torch.tensor(3., requires_grad=True)
for epoch in range(3):
    y = 3*x**2 + 4*x + 2
    y.backward()
    print(x.grad)
    x.grad.zero_()
```

Gradient Descent & Backward Pass

Computational graph and Autograd with Pytorch

The Chain Rule

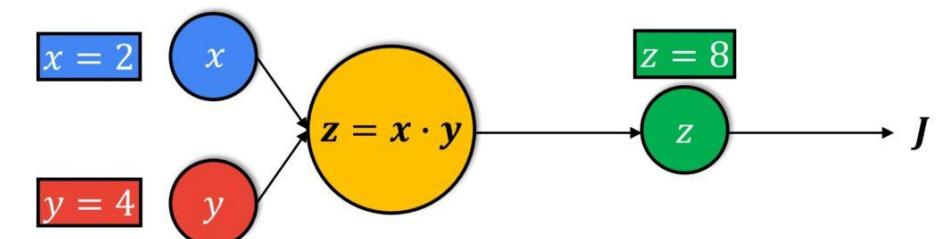
$$g = g(x) \quad f = f(x)$$



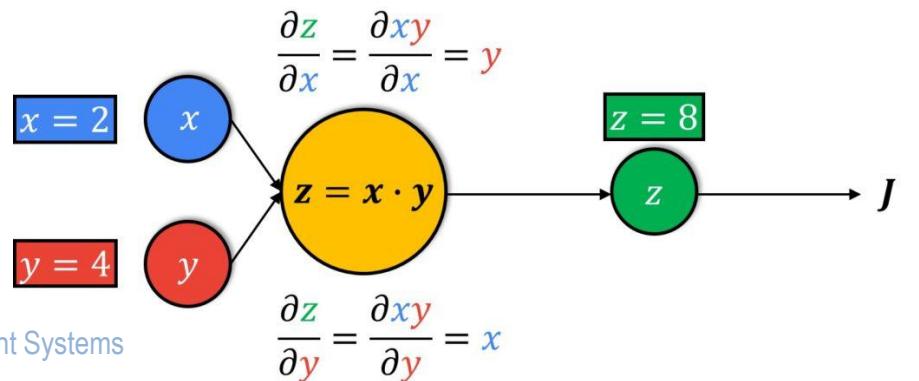
$$\frac{dg}{dx} \quad \frac{df}{dg}$$

$$\frac{d}{dx} f(g(x)) = \frac{dg}{dx} \frac{df}{dg}$$

Forward propagation



Backward propagation





Gradient Descent & Backward Pass

Computational graph and Autograd with Pytorch

4. Optimizing parameters with Autograd

Now, let's see how we can apply the chain rule in PyTorch.

```
import torch

x = torch.tensor([1.,2.,3.], requires_grad=True)
y = x * 2 + 3
z = y ** 2
print(z)
# To calculate the gradients, we would need a scalar value of z. WHY?
out = z.mean()
out.backward()

print(out)
print(x.grad)
```

Now let's see what should we do if instead of calculating the gradients for a mean value of z, we want to calculate the gradient for each z individually?

```
v = torch.tensor([1.,1.,1.])
z.backward(v)

print(x.grad/len(x))
```

Simple Neural Network

```
import numpy as np
import torch
from torch import nn
from torchvision import transforms, datasets, utils
from torch.utils.data import DataLoader, TensorDataset
```

1. Defining NN model

```
class network(torch.nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super().__init__()
        self.hidden = torch.nn.Linear(input_size, hidden_size)
        self.output = torch.nn.Linear(hidden_size, output_size)

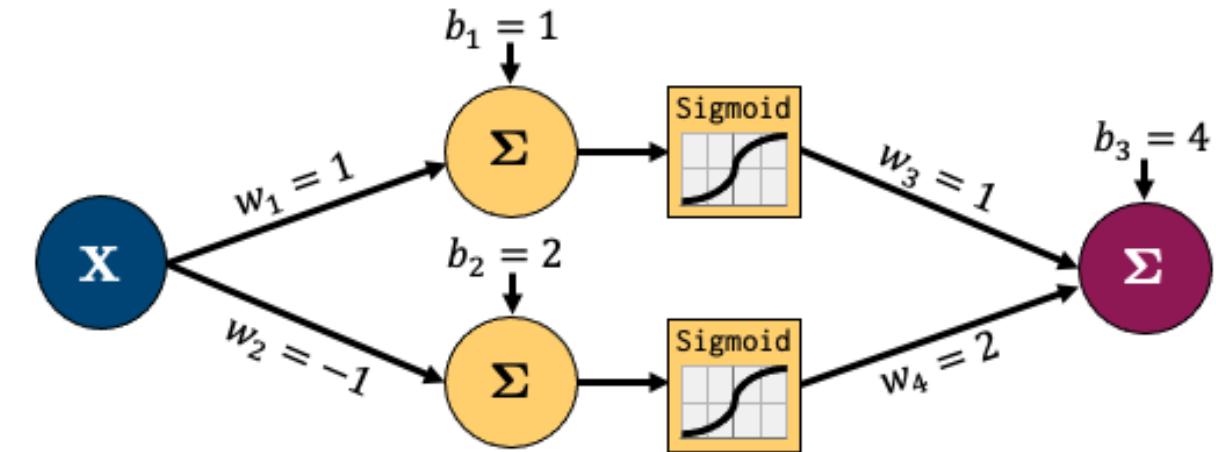
    def forward(self, x):
        x = self.hidden(x)
        x = torch.sigmoid(x)
        x = self.output(x)
        return x
```

2. Manual Weight and Bias Initialization

```
model = network(1, 2, 1) # make an instance of our network
# fix the weights manually as in the earlier figure
model.state_dict()['hidden.weight'][:] = torch.tensor([[1], [-1]])
model.state_dict()['hidden.bias'][:] = torch.tensor([1, 2])
model.state_dict()['output.weight'][:] = torch.tensor([[1, 2]])
model.state_dict()['output.bias'][:] = torch.tensor([-1])

x, y = torch.tensor([1.0]), torch.tensor([3.0]) # our x, y data

print(model.output.bias.grad)
```



3. Loss Function

```
criterion = torch.nn.MSELoss()
```

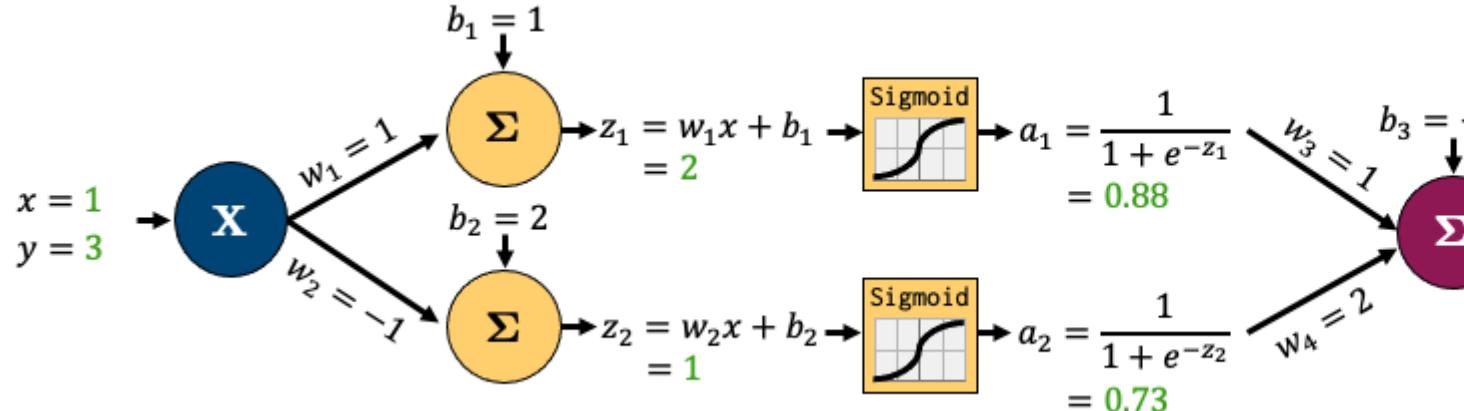
4. Forward and Backward Pass

```
loss = criterion(model(x), y)
loss.backward() # calculate gradients!
```

5. Bias Grad

```
print(model.output.bias.grad)
```

Simple Neural Network



4a. Forward Pass

$$\begin{aligned}\mathcal{L} &= \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 \\ &= (1.3429 - 3)^2 \\ &= 2.746\end{aligned}$$

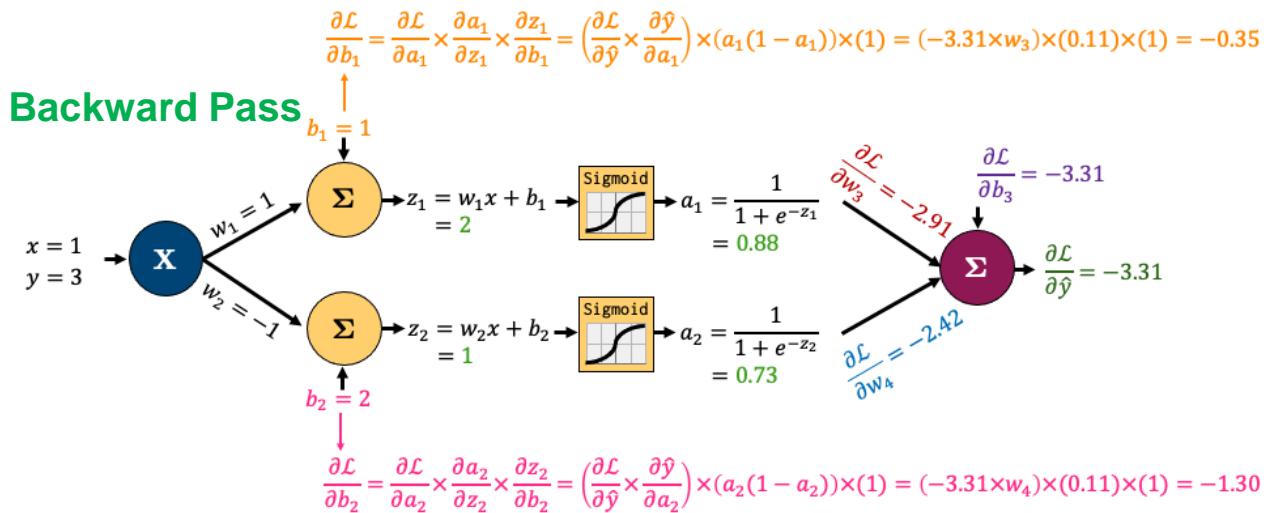
```
print("Hidden Layer Gradients")
print("Bias:", model.hidden.bias.grad)
print("Weights:", model.hidden.weight.grad.squeeze())
print()
print("Output Layer Gradients")
print("Bias:", model.output.bias.grad)
print("Weights:", model.output.weight.grad.squeeze())

model.state_dict()
```

Output

```
OrderedDict([('hidden.weight',
              tensor([[ 1.],
                     [-1.]])),
             ('hidden.bias', tensor([1., 2.])),
             ('output.weight', tensor([[1., 2.]])),
             ('output.bias', tensor([-1.]))])
```

4b. Backward Pass





Simple Neural Network

Indian Institute of Technology Kharagpur
Department of Electrical Engineering | Signal Processing and Machine Learning

6. Optimizer and Gradient Descent

```
optimizer = torch.optim.SGD(model.parameters(),  
lr=0.1)  
optimizer.step()  
  
# Model weight and bias will be different now.  
model.state_dict()
```

```
OrderedDict([('hidden.weight',  
             tensor([[ 1.0348],  
                     [-0.8697]])),  
            ('hidden.bias', tensor([1.0348, 2.1303])),  
            ('output.weight', tensor([[1.2919, 2.2423]])),  
            ('output.bias', tensor([-0.6686]))])
```

7. Weight updating over Iterations

```
optimizer.zero_grad() # Why?. Try without this command  
# After 5 iterations  
for _ in range(1, 6):  
    loss = criterion(model(x), y)  
    loss.backward()  
    print(f'b3 gradient after call {_} of loss.backward():", model.hidden.bias.grad)
```

```
b3 gradient after call 1 of loss.backward(): tensor([-0.1991, -0.5976])  
b3 gradient after call 2 of loss.backward(): tensor([-0.3983, -1.1953])  
b3 gradient after call 3 of loss.backward(): tensor([-0.5974, -1.7929])  
b3 gradient after call 4 of loss.backward(): tensor([-0.7966, -2.3906])  
b3 gradient after call 5 of loss.backward(): tensor([-0.9957, -2.9882])
```

Logistic Regression

Logistic regression is a type of regression that predicts the probability of an event.

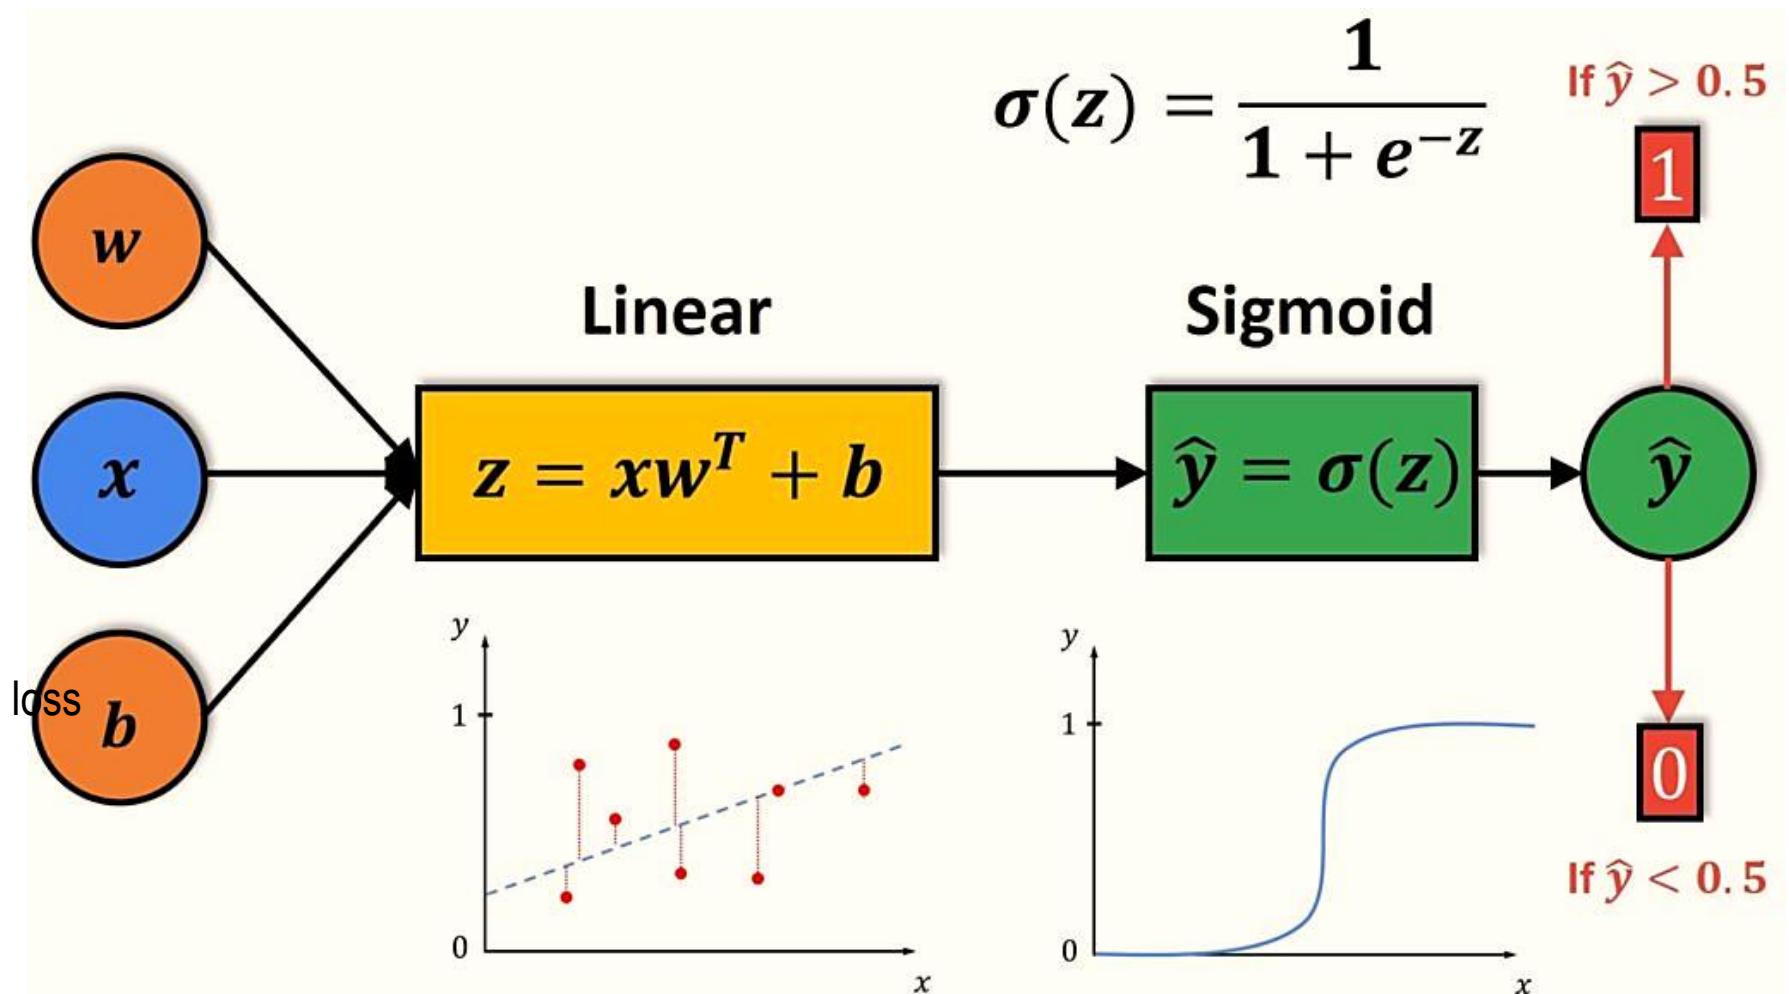
Steps for PyTorch Training

- 1) Data Preparation
- 2) Design model (input, output, forward pass with different layers)
- 3) Construct loss and optimizer
- 4) Training loop

Forward = compute prediction and loss

Backward = compute gradients

Update weights



Logistic Regression

```
import torch
import torch.nn as nn
import numpy as np
from sklearn import datasets
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
```

1) Data Preparation

```
bc = datasets.load_breast_cancer()
X, y = bc.data, bc.target

n_samples, n_features = X.shape

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=1234)

# scale
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

X_train = torch.from_numpy(X_train.astype(np.float32))
X_test = torch.from_numpy(X_test.astype(np.float32))
y_train = torch.from_numpy(y_train.astype(np.float32))
y_test = torch.from_numpy(y_test.astype(np.float32))

y_train = y_train.view(y_train.shape[0], 1)
y_test = y_test.view(y_test.shape[0], 1)
```

2) Model Definition

```
# Linear model  $f = wx + b$  , sigmoid at the end
class Model(nn.Module):
    def __init__(self, n_input_features):
        super(Model, self).__init__()
        self.linear = nn.Linear(n_input_features, 1)

    def forward(self, x):
        y_pred = torch.sigmoid(self.linear(x))
        return y_pred

model = Model(n_features)
```

3) Loss Function and Optimizer

```
num_epochs = 100
learning_rate = 0.01
criterion = nn.BCELoss()
optimizer = torch.optim.SGD(model.parameters(),
lr=learning_rate)
```

Logistic Regression

4) Training Loop

```
for epoch in range(num_epochs):
    # Forward pass and loss
    y_pred = model(X_train)
    loss = criterion(y_pred, y_train)

    # Backward pass and update
    loss.backward()
    optimizer.step()

    # zero grad before new step
    optimizer.zero_grad()

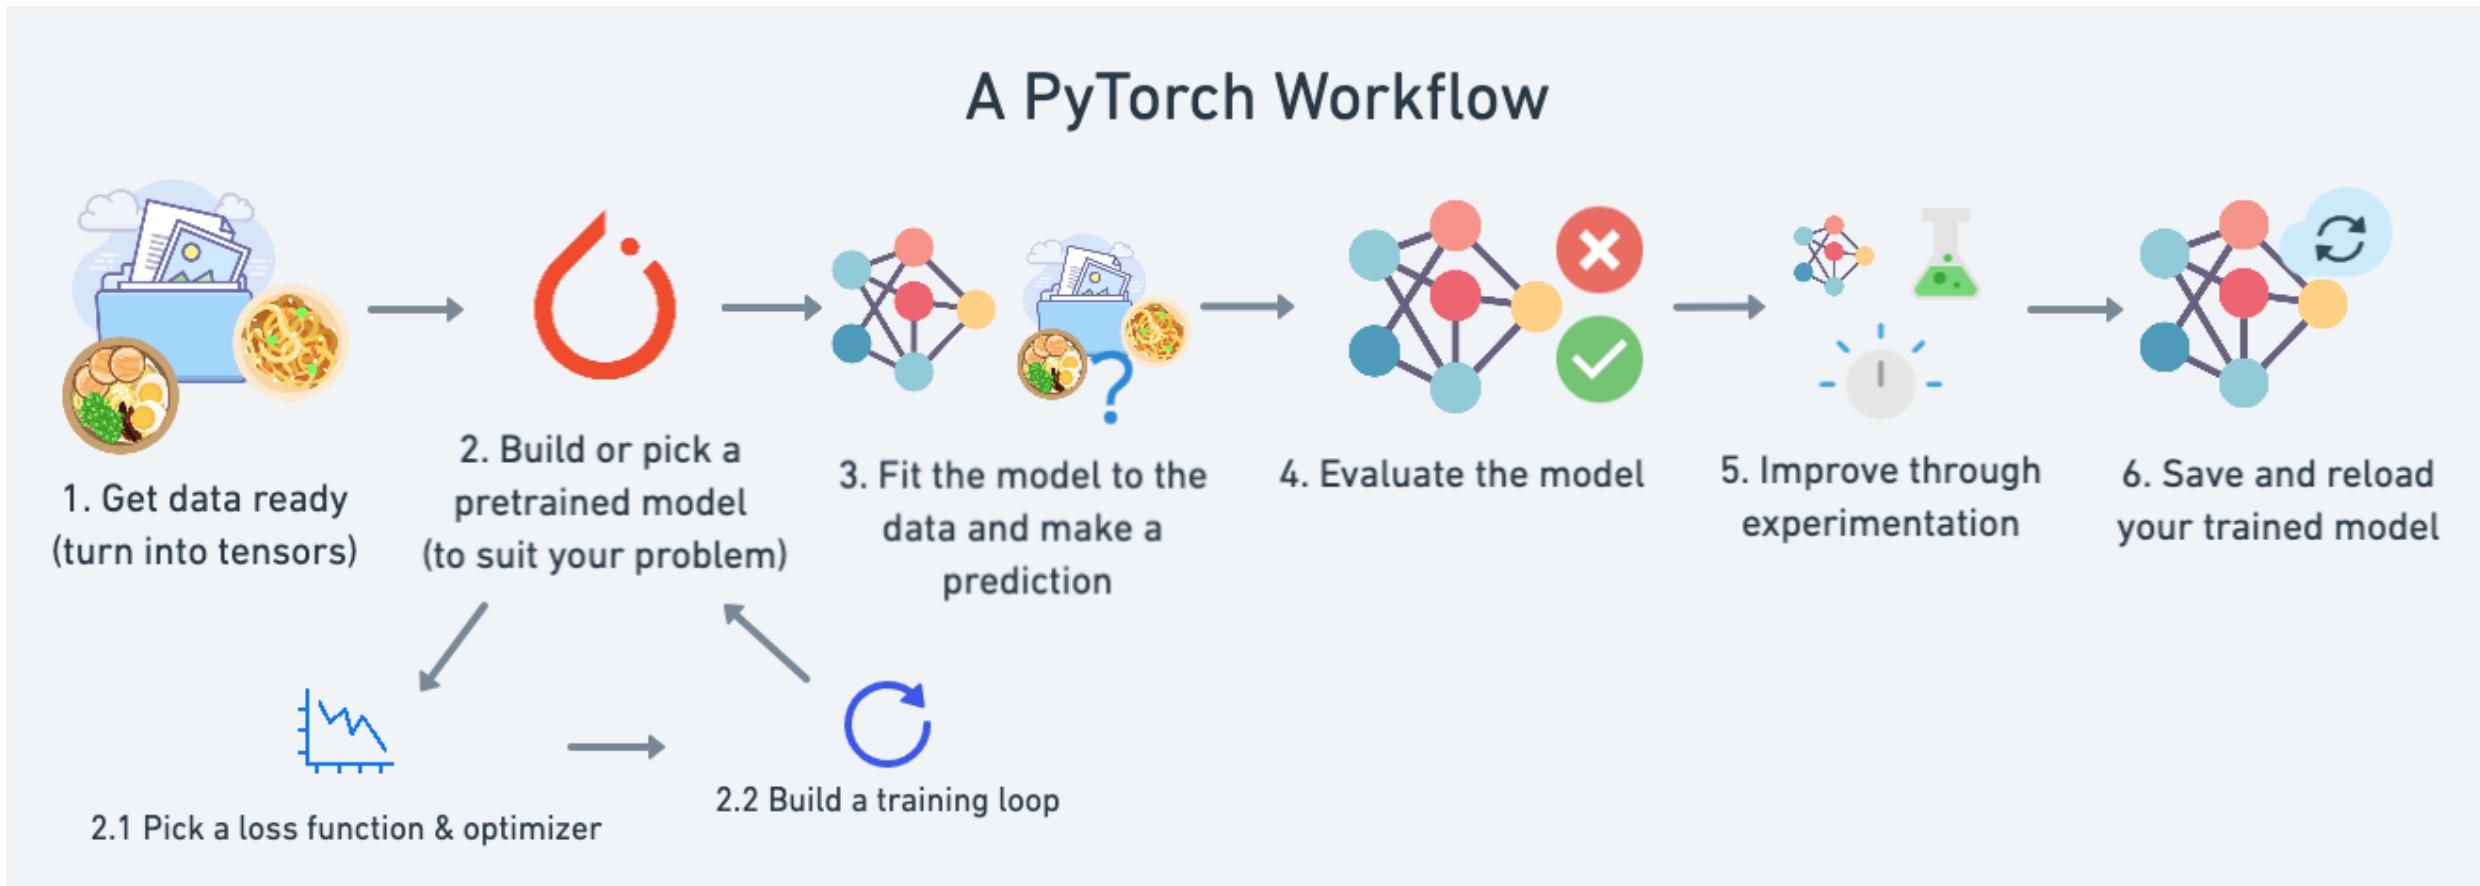
    if (epoch+1) % 10 == 0:
        print(f'epoch: {epoch+1}, loss = {loss.item():.4f}')
```

```
with torch.no_grad():
    y_predicted = model(X_test)
    y_predicted_cls = y_predicted.round()
    acc = y_predicted_cls.eq(y_test).sum() / float(y_test.shape[0])
    print(f'accuracy: {acc.item():.4f}')
```

Testing and Performance Analysis

PyTorch Workflow for DL

Various Class Definitions required for Building ANN/CNN Models



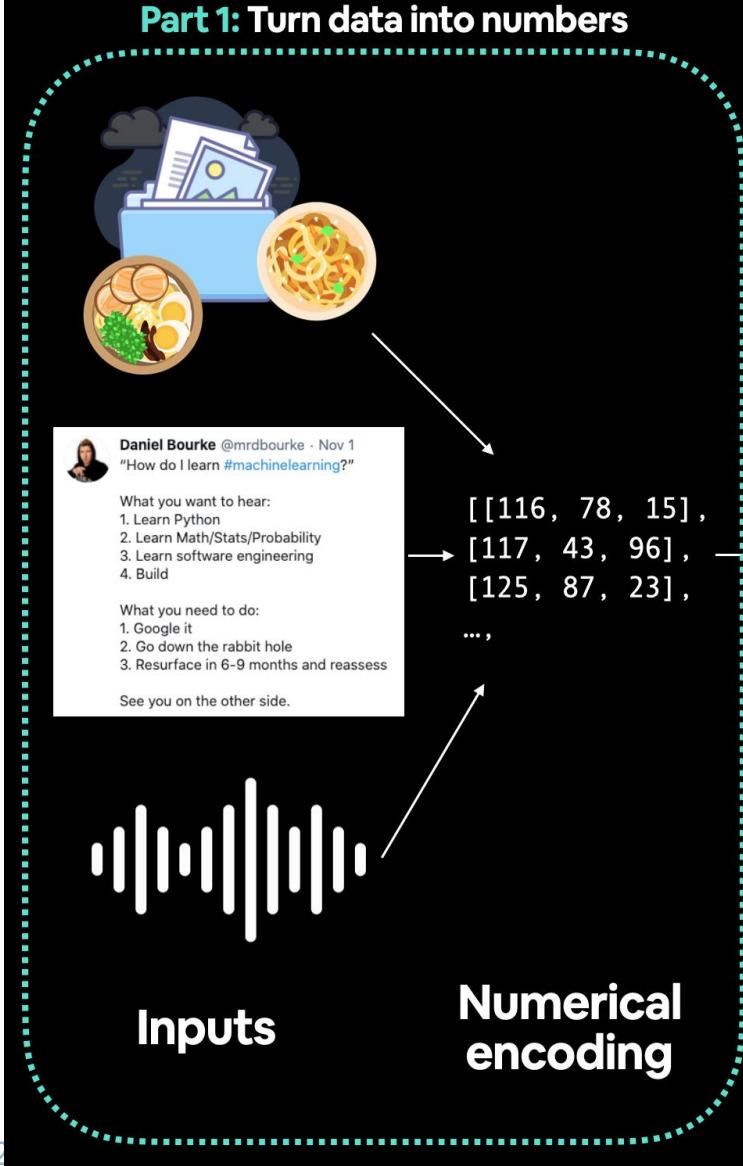
<https://www.learnpytorch.io/>

PyTorch Workflow for DL

Create DataLoader & Data-Transformer

Indian Institute of Technology Kharagpur
Department of Electrical Engineering | Signal Processing and Machine Learning

Part 1: Turn data into numbers



```
import torch
from torch.utils.data import Dataset, DataLoader
import numpy as np
```

```
class CustomDataset(Dataset):
    def __init__(self, data, labels):
        """
        Args:
            data (numpy array or torch.Tensor): Input features.
            labels (numpy array or torch.Tensor): Corresponding labels.
        """

```

Args:

data (numpy array or torch.Tensor): Input features.
labels (numpy array or torch.Tensor): Corresponding labels.
"""

```
self.data = torch.tensor(data, dtype=torch.float32)
self.labels = torch.tensor(labels, dtype=torch.long)
```

```
def __len__(self):
```

"""Returns the total number of samples."""
return len(self.data)

```
def __getitem__(self, index):
```

"""Retrieves a sample at the given index."""
return self.data[index], self.labels[index]



PyTorch Workflow for DL

Create DataLoader & Data-Transformer

Indian Institute of Technology Kharagpur
Department of Electrical Engineering | Signal Processing and Machine Learning

| Feature | PyTorch DataLoader | PyTorch Transformers (Transforms) |
|------------------------------|---|---|
| Purpose | Efficiently loads and batches dataset | Applies data augmentation and preprocessing |
| Class Location | <code>torch.utils.data.DataLoader</code> | <code>torchvision.transforms</code> (for images) |
| Primary Function | Iterates through dataset with batching, shuffling, and parallel processing | Modifies dataset samples (e.g., resizing, normalization) |
| Usage | Used in training loops to load data efficiently | Applied to each sample before it is returned by the dataset |
| Works With | Any <code>Dataset</code> object | Image datasets (<code>torchvision</code>), tabular/text data |
| Key Parameters | <code>batch_size</code> , <code>shuffle</code> , <code>num_workers</code> , <code>drop_last</code> | <code>Resize</code> , <code>Normalize</code> , <code>RandomCrop</code> , <code>ToTensor</code> , etc. |
| Example Usage | <code>dataloader = DataLoader(dataset, batch_size=32, shuffle=True)</code> | <code>transform = transforms.Compose([transforms.Resize(256), transforms.ToTensor()])</code> |
| Supports Custom Datasets? | <input checked="" type="checkbox"/> Yes, via <code>torch.utils.data.Dataset</code> | <input checked="" type="checkbox"/> Yes, via <code>__getitem__</code> in custom datasets |
| Parallel Processing? | <input checked="" type="checkbox"/> Yes (<code>num_workers > 0</code> speeds up loading) | <input type="checkbox"/> No, applied within dataset class |
| Application in Deep Learning | Handles mini-batches, making training efficient | Prepares data for better model generalization |

```
import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transforms
import matplotlib.pyplot as plt

# Device configuration
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

# MNIST dataset
train_dataset = torchvision.datasets.MNIST(root='./data',
                                             train=True,
                                             transform=transforms.ToTensor(),
                                             download=True)

test_dataset = torchvision.datasets.MNIST(root='./data',
                                         train=False,
                                         transform=transforms.ToTensor())

# Data loader
train_loader = torch.utils.data.DataLoader(dataset=train_dataset,
                                            batch_size=100,
                                            shuffle=True)

test_loader = torch.utils.data.DataLoader(dataset=test_dataset,
                                         batch_size=batch_size,
                                         shuffle=False)
```



PyTorch Workflow for DL

Creating custom DataLoader

Indian Institute of Technology Kharagpur
Department of Electrical Engineering | Signal Processing and Machine Learning

```
import torch
from torch.utils.data import Dataset, DataLoader
import torchvision.transforms as transforms
import torchvision.datasets as datasets
import matplotlib.pyplot as plt
import numpy as np
```

Device configuration

```
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

- **ToTensor()**: Converts PIL images to PyTorch tensors.
- **Normalize()**: Scales pixel values to range [-1,1][−1,1] for better convergence.

```
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))
])
```

Load training dataset with custom dataset class

```
train_dataset = CustomMNISTDataset(train=True, transform=transform)
train_loader = DataLoader(train_dataset, batch_size=8, shuffle=True)
```

```
class CustomMNISTDataset(Dataset):
    def __init__(self, train=True, transform=None):
        self.mnist_data = datasets.MNIST(root='./data', train=train, download=True,
                                         transform=transform)

    def __len__(self):
        return len(self.mnist_data)

    def __getitem__(self, idx):
        image, label = self.mnist_data[idx]
        return image, label
```

```
# Get a batch of data
data_iter = iter(train_loader)
images, labels = next(data_iter)

# Function to display images
def imshow(img):
    img = img / 2 + 0.5 # Unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)), cmap="gray")
    plt.axis("off")
    plt.show()

# Display batch of images
imshow(torchvision.utils.make_grid(images))
print("Labels:", labels.numpy())
```

PyTorch Workflow for DL

DataLoader – Splitting Dataset

| Split | Purpose | Amount of total data | How often is it used? |
|----------------|--|----------------------|-----------------------|
| Training set | The model learns from this data (like the course materials you study during the semester). | ~60-80% | Always |
| Validation set | The model gets tuned on this data (like the practice exam you take before the final exam). | ~10-20% | Often but not always |
| Testing set | The model gets evaluated on this data to test what it has learned (like the final exam you take at the end of the semester). | ~10-20% | Always |

1. Training Set (80%)

- Used to train the model by adjusting weights through backpropagation.
- The model learns patterns and representations from this dataset.

2. Validation Set (10%)

- Used to fine-tune hyperparameters (e.g., learning rate, batch size, dropout rate).
- Helps in preventing overfitting, ensuring the model does not memorize the training data.
- Early stopping is often based on validation loss improvement.

3. Test Set (10%)

- Used only after the model is finalized to evaluate its performance.
- Measures how well the model generalizes to completely unseen data.
- Ensures unbiased model assessment before deployment.



PyTorch Workflow for DL

DataLoader – Splitting Dataset

```
from torch.utils.data import random_split

# Define dataset
full_dataset = CustomMNISTDataset(train=True, transform=transform)

total_size = len(full_dataset)
train_size = int(0.8 * total_size)
val_size = int(0.1 * total_size)
test_size = total_size - train_size - val_size # Ensures no data loss

# Split dataset
train_dataset, val_dataset, test_dataset = random_split(full_dataset, [train_size, val_size, test_size])

# Create DataLoaders
train_loader = DataLoader(train_dataset, batch_size=100, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=100, shuffle=False)
test_loader = DataLoader(test_dataset, batch_size=100, shuffle=False)
```

PyTorch Workflow for DL

Model Definition

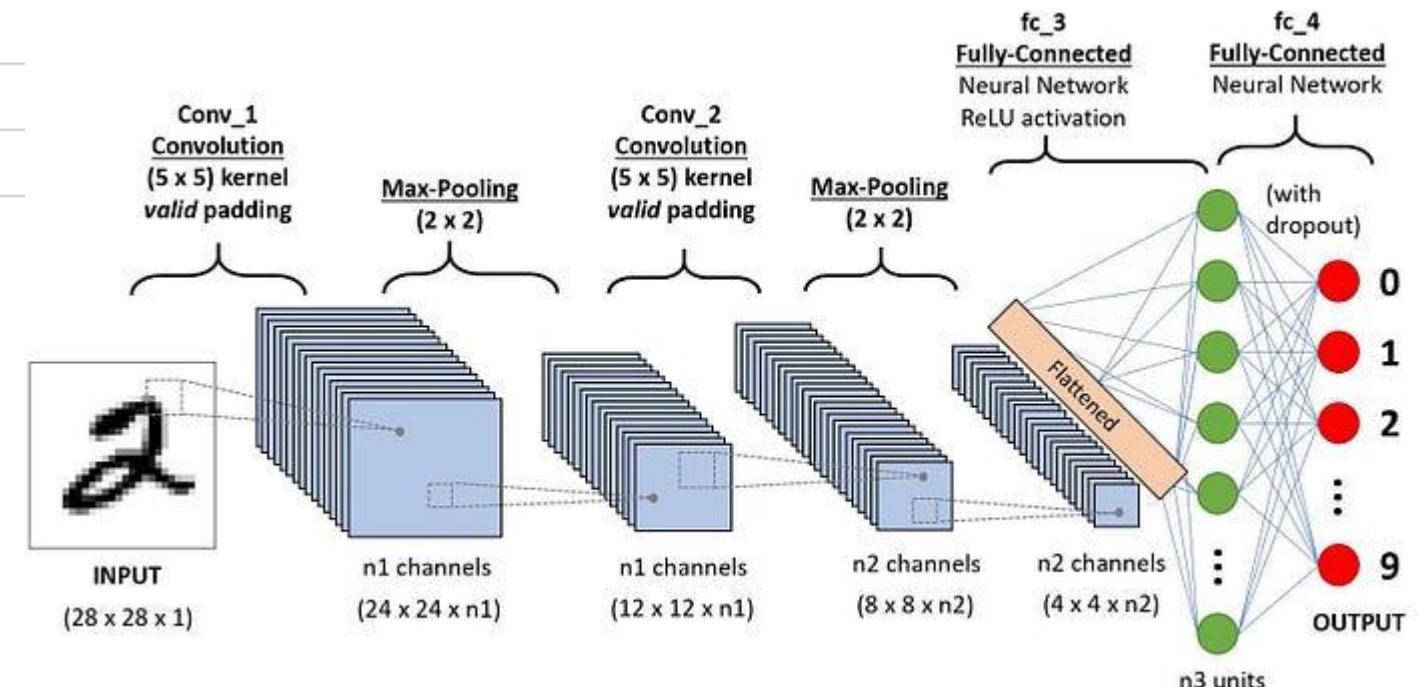
Architecture Summary

| Layer | Details |
|-----------------------------|--|
| Input | $28 \times 28 \times 1$ |
| Conv2D (1st layer) | 5 × 5 filter, 32 filters, ReLU, MaxPooling (2×2) |
| Conv2D (2nd layer) | 5 × 5 filter, 64 filters, ReLU, MaxPooling (2×2) |
| Flatten 1 | Converts feature maps into a 1D vector |
| Fully Connected Layer (fc1) | 256 neurons, ReLU Activation |
| Dropout | Dropout probability 0.5 |
| Fully Connected Layer (fc2) | 10-class output |

```

import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
import torch.nn.functional as F

from torch.utils.data import DataLoader
    
```



PyTorch Workflow for DL

Model Definition

```

class CustomCNN(nn.Module):
    def __init__(self, num_classes=10):
        super(CustomCNN, self).__init__()

        # First Conv Layer: 1 input channel, 32 filters, kernel size 5x5
        self.conv1 = nn.Conv2d(in_channels=1, out_channels=32, kernel_size=5)
        self.pool = nn.MaxPool2d(kernel_size=2, stride=2)

        # Second Conv Layer: 32 input channels, 64 filters, kernel size 5x5
        self.conv2 = nn.Conv2d(in_channels=32, out_channels=64, kernel_size=5)

        # Fully Connected Layers
        self.fc1 = nn.Linear(64 * 4 * 4, 256) # Adjusted for input size after convolutions
        self.fc2 = nn.Linear(256, num_classes)

        # Dropout
        self.dropout = nn.Dropout(0.5)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x))) # Conv1 -> ReLU -> MaxPool
        x = self.pool(F.relu(self.conv2(x))) # Conv2 -> ReLU -> MaxPool

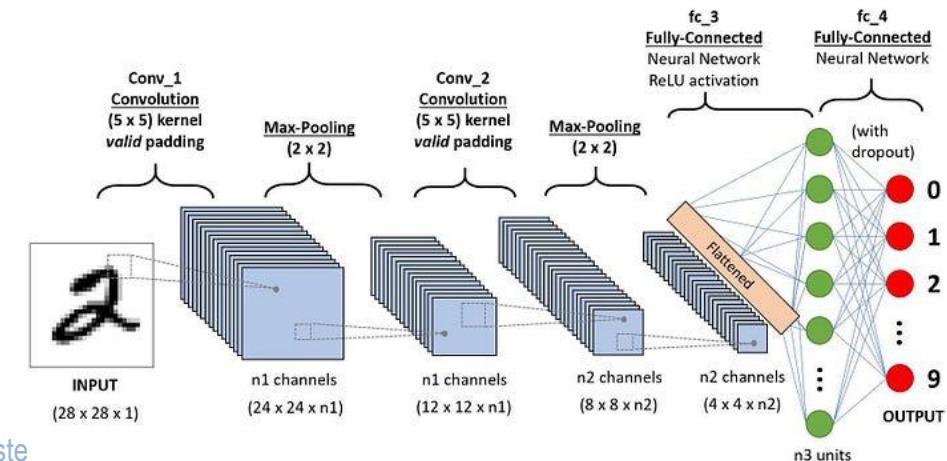
        x = torch.flatten(x, start_dim=1) # Flattening layer (Flatten1)
        x = F.relu(self.fc1(x)) # Fully connected layer with ReLU activation
        x = self.dropout(x) # Dropout layer

        x = self.fc2(x) # Final output layer
        return x

```

The base class for all neural network modules, all the building blocks for neural networks are subclasses. If you're building a neural network in PyTorch, your models should subclass `nn.Module`

All `nn.Module` subclasses require a `forward()` method, this defines the computation that will take place on the data passed to the particular `nn.Module`.





PyTorch Workflow for DL

Training Loop

| Step name | What does it do? | Code example |
|---|---|--|
| Forward pass | The model goes through all of the training data once, performing its <code>forward()</code> function calculations. | <code>model(x_train)</code> |
| Calculate the loss | The model's outputs (predictions) are compared to the ground truth and evaluated to see how wrong they are. | <code>loss = loss_fn(y_pred, y_train)</code> |
| Zero gradients | The optimizers gradients are set to zero (they are accumulated by default) so they can be recalculated for the specific training step. | <code>optimizer.zero_grad()</code> |
| Perform backpropagation on the loss | Computes the gradient of the loss with respect for every model parameter to be updated (each parameter with <code>requires_grad=True</code>). This is known as backpropagation, hence "backwards". | <code>loss.backward()</code> |
| Update the optimizer (gradient descent) | Update the parameters with <code>requires_grad=True</code> with respect to the loss gradients in order to improve them. | <code>optimizer.step()</code> |

```
# Initialize model
model = CustomCNN()

# Define optimizer
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Define loss function
criterion = nn.CrossEntropyLoss()
```



Epoch loop

Mini-batch loop

zero gradients `model.zero_grad()`

forward pass `outputs=model(inputs)`

compute loss `loss=criterion(outputs, labels)`

compute gradients `loss.backward()`

update weights `optimizer.step()`

Training Loop

Training loop

```
for epoch in range(5):
    model.train()
    running_loss = 0.0

    for data in trainloader:

        inputs, labels = data
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item()

    accuracy = evaluate(model)

    print(f"Epoch {epoch + 1},
          Loss: {running_loss / len(trainloader)},
          Accuracy: {accuracy * 100:.2f}%
          "")
```



PyTorch Workflow for DL

Model Evaluation

```
def evaluate(model):  
  
    model.eval()  
    correct, total = 0, 0  
  
    with torch.no_grad():  
        for data in testloader:  
            inputs, labels = data  
            outputs = model(inputs)  
            _, predicted = torch.max(outputs.data, 1)  
            total += labels.size(0)  
            correct += (predicted == labels).sum().item()  
  
    return correct / total
```

Evaluate model

Step 1) Import required packages and libraries

```
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
import torch.nn.functional as F

from torch.utils.data import DataLoader
```

Step 3) Define the PyTorch Model
Architecture

```
class PyTorchNet(nn.Module):

    def __init__(self):
        super(PyTorchNet, self).__init__()
        self.fc1 = nn.Linear(28*28, 512)
        self.fc2 = nn.Linear(512, 1024)
        self.fc3 = nn.Linear(1024, 128)
        self.fc4 = nn.Linear(128, 10)

    def forward(self, x):
        x = x.view(-1, 28*28)
        x = torch.relu(self.fc1(x))
        x = torch.relu(self.fc2(x))
        x = torch.relu(self.fc3(x))
        x = self.fc4(x)
        return x
```

Forward pass

Step 2) Load the dataset

```
# 1) Data Transformer
transform = transforms.Compose([transforms.ToTensor(),
                               transforms.Normalize((0.5,), (0.5,))])

# 2) Create Train Dataset
trainset = torchvision.datasets.MNIST(root='./data', train=True,
                                       download=True, transform=transform)
trainloader = DataLoader(trainset, batch_size=64, shuffle=True)

# 3) Create Test Dataset
testset = torchvision.datasets.MNIST(root='./data', train=False,
                                       download=True, transform=transform)
testloader = DataLoader(testset, batch_size=64, shuffle=False)
```

Step 4) Initialize the model and define the loss function and optimizer

```
# Initialize model
model = PyTorchNet()

# Define optimizer
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Define loss function
criterion = nn.CrossEntropyLoss()
```

```
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
import torch.nn.functional as F

from torch.utils.data import DataLoader
```

```
class PyTorchNet(nn.Module):

    def __init__(self):
        super(PyTorchNet, self).__init__()
        self.fc1 = nn.Linear(28*28, 512)
        self.fc2 = nn.Linear(512, 1024)
        self.fc3 = nn.Linear(1024, 128)
        self.fc4 = nn.Linear(128, 10)

    def forward(self, x):
        x = x.view(-1, 28*28)
        x = torch.relu(self.fc1(x))
        x = torch.relu(self.fc2(x))
        x = torch.relu(self.fc3(x))
        x = self.fc4(x)
        return x
```

```
# 1) Data Transformation
transform = transforms.Compose([
    transforms.RandomResizedCrop(224),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.45, 0.406], [0.229, 0.224, 0.225])
])

# 2) Create Train and Test Datasets
trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                        download=True, transform=transform)
testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                       download=True, transform=transform)

# 3) Create Train and Test DataLoaders
trainloader = DataLoader(trainset, batch_size=4,
                        shuffle=True, num_workers=2)
testloader = DataLoader(testset, batch_size=4,
                       shuffle=False, num_workers=2)

# Initialize Model
model = PyTorchNet()

# Define optimizer
optimizer = optim.SGD(model.parameters(), lr=0.001)

# Define loss function
criterion = nn.CrossEntropyLoss()

# Training loop
for epoch in range(5):
    model.train()
    running_loss = 0.0

    for data in trainloader:
        inputs, labels = data
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item()

    accuracy = evaluate(model)
    print(f"Epoch {epoch + 1},\nLoss: {running_loss / len(trainloader)},\nAccuracy: {accuracy * 100:.2f}%")
```

Training loop

Step 5) Train the model



CNN Complete Model

Indian Institute of Technology Kharagpur
Department of Electrical Engineering | Signal Processing and Machine Learning

```
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
import torch.nn.functional as F

from torch.utils.data import Dataset, DataLoader

class PyTorchNet(nn.Module):

    def __init__(self):
        super(PyTorchNet, self)
        self.fc1 = nn.Linear(28*28, 128)
        self.fc2 = nn.Linear(128, 128)
        self.fc3 = nn.Linear(128, 128)
        self.fc4 = nn.Linear(128, 10)

    def forward(self, x):
        x = x.view(-1, 28*28)
        x = torch.relu(self.fc1(x))
        x = torch.relu(self.fc2(x))
        x = torch.relu(self.fc3(x))
        x = self.fc4(x)
        return x
```

Evaluate model

Training loop

```
def evaluate(model):

    model.eval()
    correct, total = 0, 0

    with torch.no_grad():
        for data in testloader:
            inputs, labels = data
            outputs = model(inputs)
            _, predicted = torch.max(outputs, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

    return correct / total
```

Step 6) Define the evaluation method

```
) , labels)
ax(outputs.data, 1)
)
= labels.sum().item() n()

print(f"""Epoch {epoch + 1},
      Loss: {running_loss / len(trainloader)},
      Accuracy: {accuracy * 100:.2f}%
      """)
```

PyTorch Workflow for DL

Model Inference

| PyTorch method | What does it do? |
|--|---|
| <code>torch.save</code> | Saves Model, tensors and various other Python objects like dictionaries can be saved using <code>torch.save</code> . |
| <code>torch.load</code> | Loads Python object files (like models, tensors or dictionaries) into memory. You can also set which device to load the object to (CPU, GPU etc). |
| <code>torch.nn.Module.load_state_dict</code> | Loads a model's parameter dictionary (<code>model.state_dict()</code>) using a saved <code>state_dict()</code> object. |

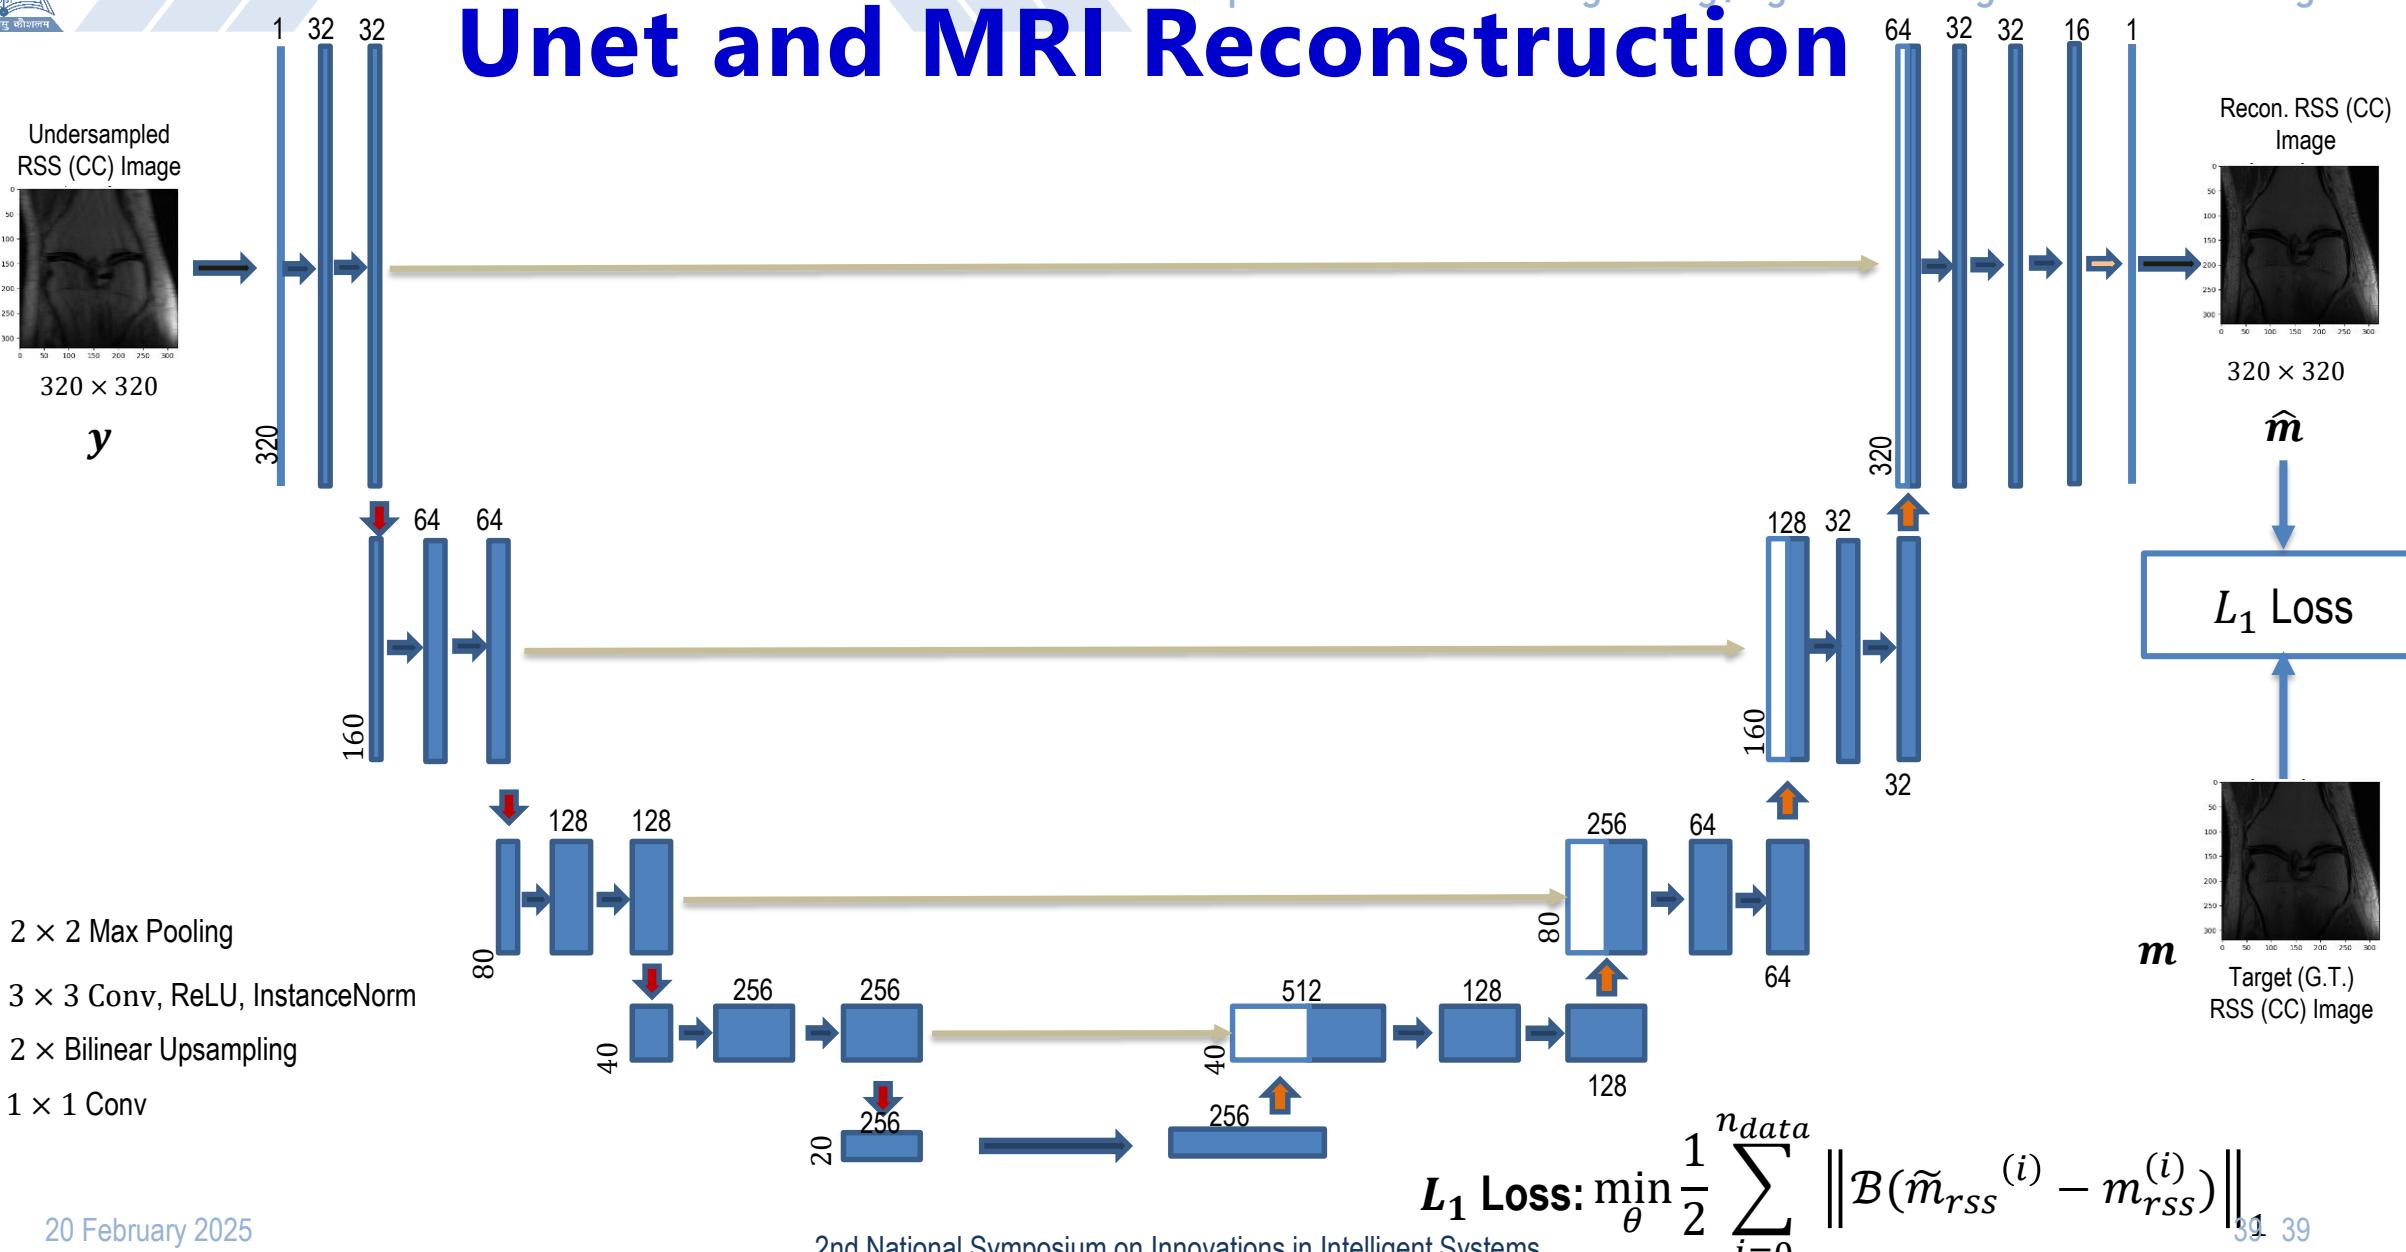
```
# Save the trained model
torch.save(model.state_dict(), 'model.pth')
```

```
# Load the model
model_loaded = PyTorchNet()
model_loaded.load_state_dict(torch.load('model.pth'))
model_loaded.eval()
```

```
# Inference: Example of prediction
with torch.no_grad():
    for data in testloader:
        inputs, labels = data
        outputs = model_loaded(inputs)
        _, predicted = torch.max(outputs, 1)
        print(f"Predicted: {predicted}, Ground Truth: {labels}")
        break # Just use one batch for demonstration
```

```
# Evaluate the loaded model
accuracy = evaluate(model_loaded)
print(f"Loaded model accuracy: {accuracy * 100:.2f}%")
```

Unet and MRI Reconstruction



Model Complexity Analysis

1. torchsummary

```
pip install torchsummary
```

```
import torch
import torch.nn as nn
import torch.nn.functional as F
from torchsummary import summary

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
        self.conv2_drop = nn.Dropout2d()
        self.fc1 = nn.Linear(320, 50)
        self.fc2 = nn.Linear(50, 10)

    def forward(self, x):
        x = F.relu(F.max_pool2d(self.conv1(x), 2))
        x = F.relu(F.max_pool2d(self.conv2_drop(self.conv2(x)), 2))
        x = x.view(-1, 320)
        x = F.relu(self.fc1(x))
        x = F.dropout(x, training=self.training)
        x = self.fc2(x)
        return F.log_softmax(x, dim=1)

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = Net().to(device)

summary(model, (1, 28, 28))
```

| Layer (type) | Output Shape | Param # |
|--------------|-------------------|---------|
| Conv2d-1 | [-1, 10, 24, 24] | 260 |
| Conv2d-2 | [-1, 20, 8, 8] | 5,020 |
| Dropout2d-3 | [-1, 20, 8, 8] | 0 |
| Linear-4 | [-1, 50] | 16,050 |
| Linear-5 | [-1, 10] | 510 |

Total params: 21,840

Trainable params: 21,840

Non-trainable params: 0

Input size (MB): 0.00

Forward/backward pass size (MB): 0.06

Params size (MB): 0.08

Estimated Total Size (MB): 0.15

2. GFLOP and GMAC

pip install ptflops

$$GMAC = \frac{(Number\ of\ Multiply - Add\ Operations)}{10^9}$$

```
import torchvision.models as models
import torch
from ptflops import get_model_complexity_info
import re

#Model that's already available
net = models.densenet161()
macs, params = get_model_complexity_info(net, (3, 224, 224), as_strings=True,
print_per_layer_stat=True, verbose=True)
# Extract the numerical value
flops = eval(re.findall(r'([\d.]+)', macs)[0])*2
# Extract the unit
flops_unit = re.findall(r'([A-Za-z]+)', macs)[0][0]

print('Computational complexity: {:.8}'.format(macs))
print('Computational complexity: {} {}Flops'.format(flops, flops_unit))
print('Number of parameters: {:.8}'.format(params))
```

$$1\ GFLOP = 2 * GMAC$$

$$GFLOP = \frac{(Number\ of\ Floating - Point\ Operations)}{\frac{(Elapsed\ Time\ in\ Seconds)}{(10^9)}}$$

Computational complexity: 7.82 GMac
 Computational complexity: 15.64 GFlops
 Number of parameters: 28.68 M

<https://medium.com/@ajithkumary/how-to-calculate-gmac-and-gflops-with-python-62004a753e3b>



Model Complexity Analysis

<https://pytorch.org/blog/zeus/>

3. Deep Learning Energy Measurement

```
pip install zeus-ml==0.7.0
```

```
from zeus.monitor import ZeusMonitor

# All four GPUs are measured simultaneously.
monitor = ZeusMonitor(gpu_indices=[0,1,2,3])

# Measure total time and energy within the window.
monitor.begin_window("training")

for e in range(100):
    # Measurement windows can arbitrarily be overlapped.
    monitor.begin_window("epoch")
    for x, y in train_dataloader:
        y_hat = model(x)
        loss = criterion(y, y_hat)
        loss.backward()
        optim.step()
    measurement = monitor.end_window("epoch")
    print(f"Epoch {e}: {measurement.time} s, {measurement.total_energy} J")

measurement = monitor.end_window("training")
print(f"Entire training: {measurement.time} s, {measurement.total_energy} J")
```



4. CodeCarbon

pip install codecarbon

Estimate and

Estimates your
(GPU + CPU + P
intensity of the r

```
from codecarbon import *

@track_emissions
def training_loop():
    # Compute intensive
```



Thank you for your time

and attention ...

Thanks to the organizers of **National Symposium – 2025**

Questions?

A Special Thanks to my Mentor:
Dr. Debdoott Sheet

Associate Professor, IIT Kgp.



<https://www.linkedin.com/in/susant-panigrahi-b79b1396/>



<https://scholar.google.co.in/citations?user=1MjsFuUAAAAJ&hl=en>



<https://github.com/susant146?tab=repositories>

It's all happening here:

- Kharagpur Learning, Imaging & Visualization (**KLIV**) research group: <https://iitkliv.github.io/>